System overview

One request path, one observation path, and an approved response gate.

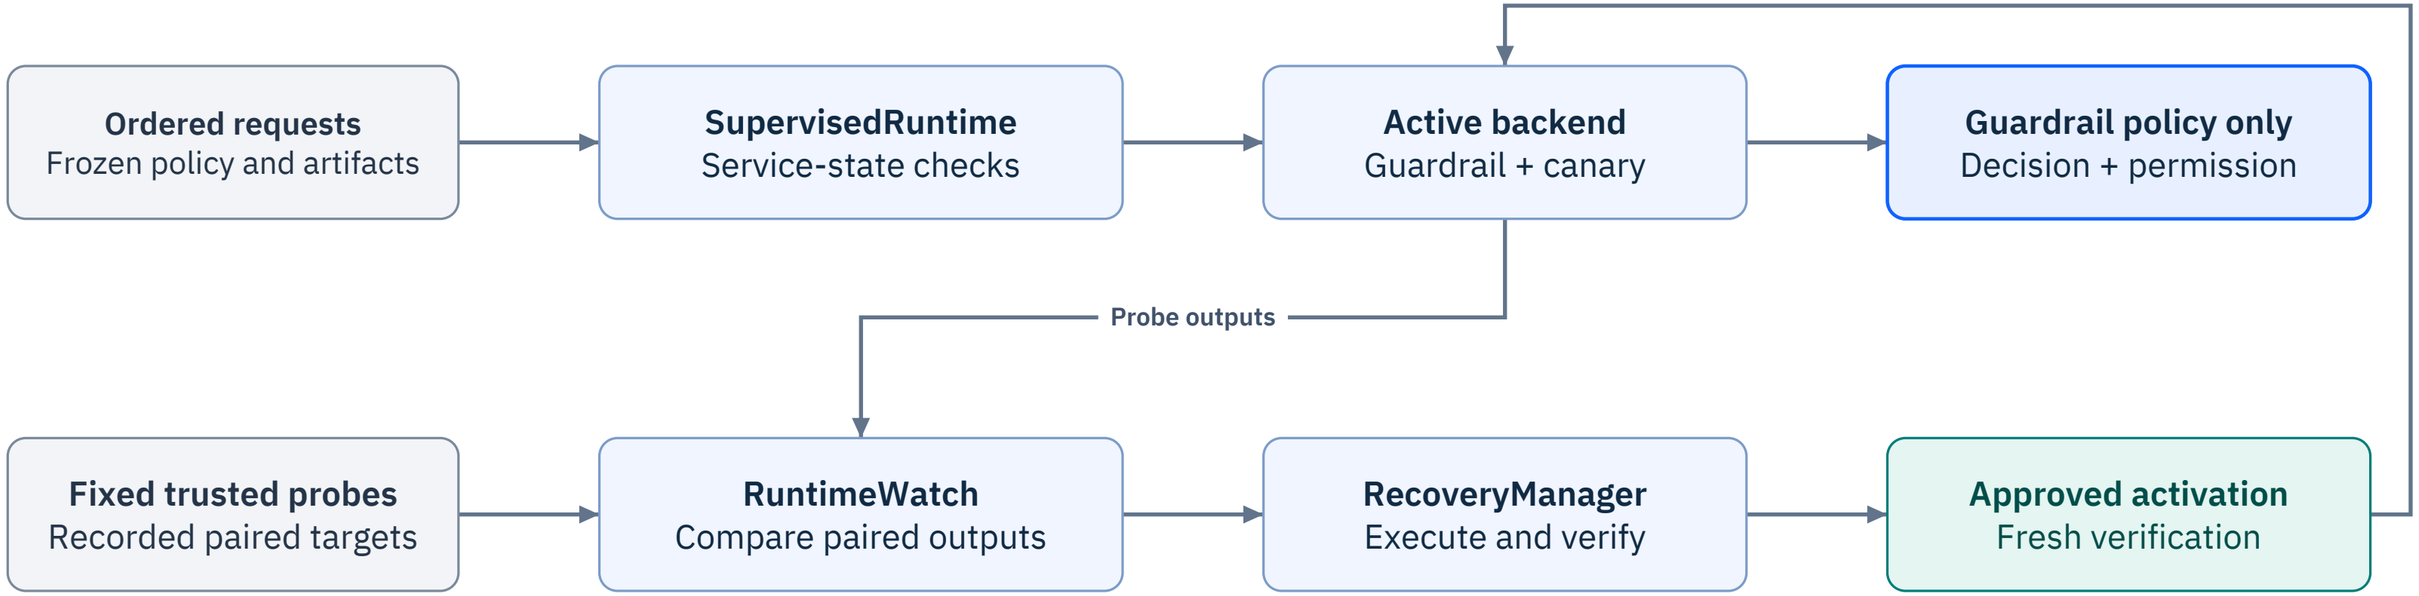

The guardrail decides. The canary observes. Both requests and probes use the active backend.

The supervisor returns a decision and an execution-permission flag. The selected replay records the flag; it executes no customer action.

Request decision flow

Only the guardrail policy can return allow or refuse.

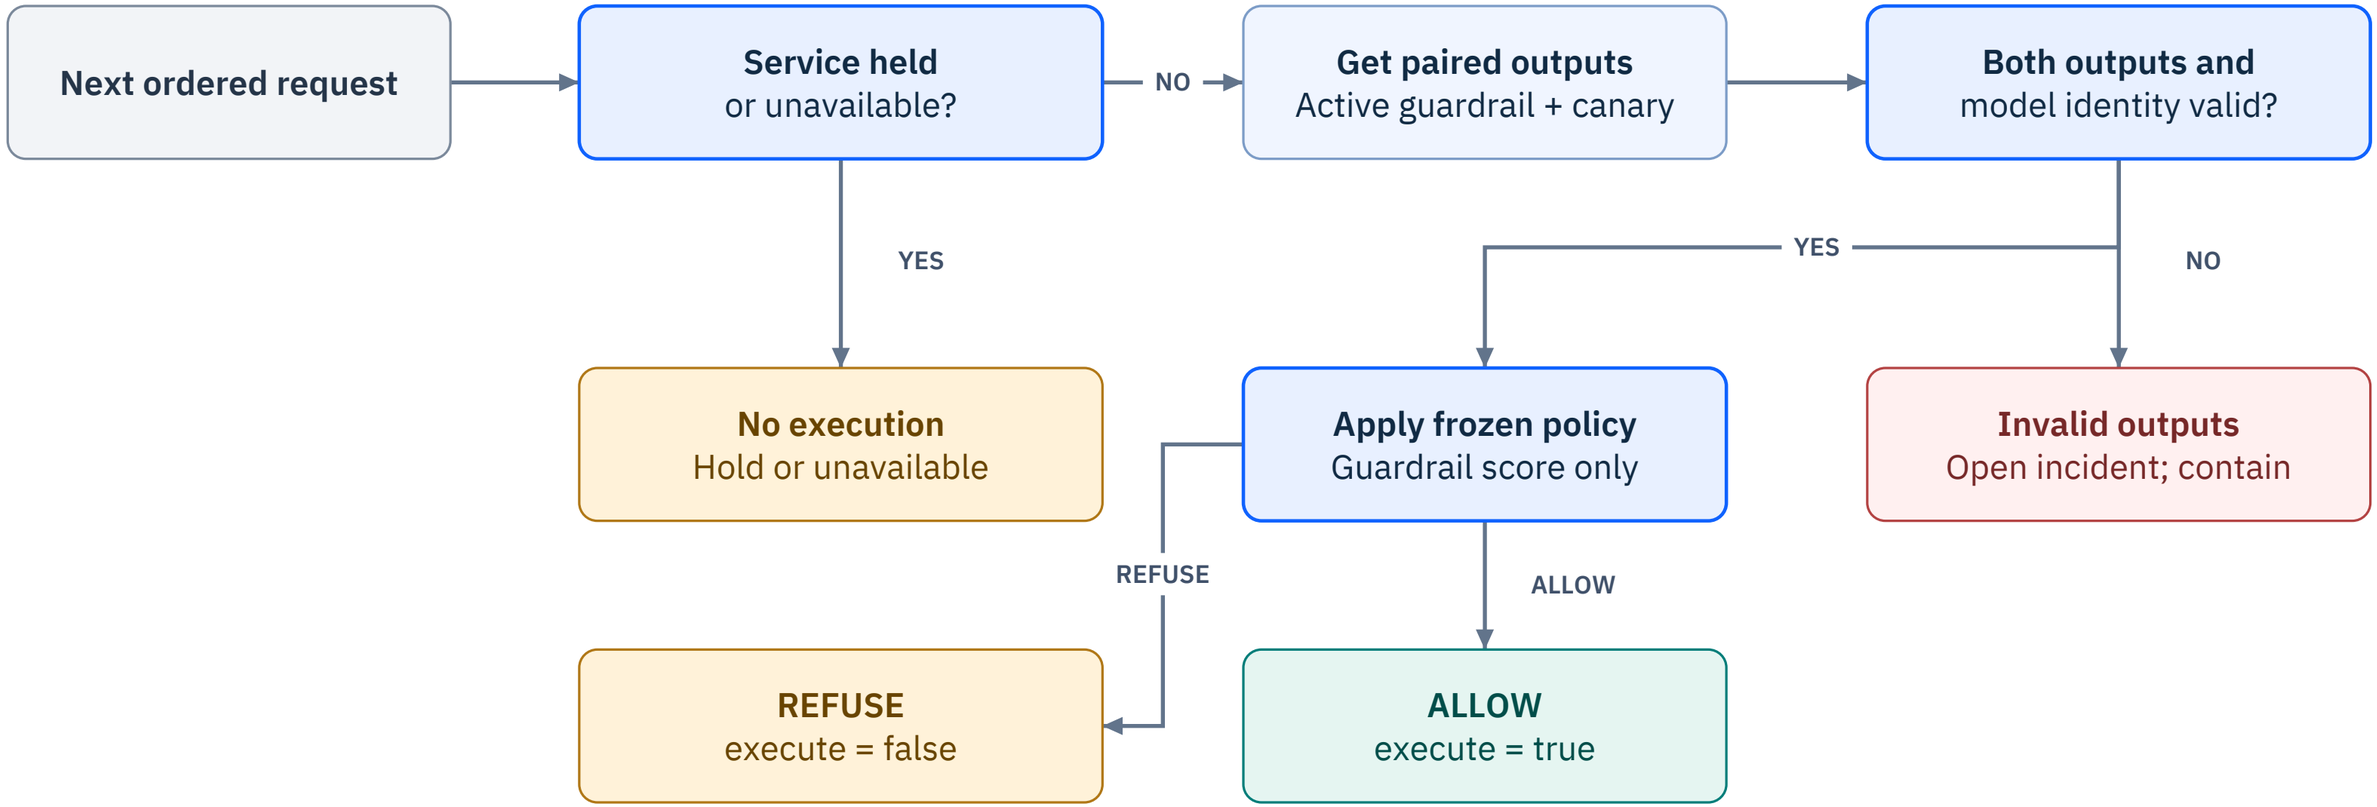

A hold, unavailable service, invalid outputs, or refusal blocks execution. These outcomes retain their separate decision types.
An allow result permits a caller action; it does not prove that the request is universally safe.

Fresh verification before activation

A manager-owned, single-use plan cannot self-certify its candidate.

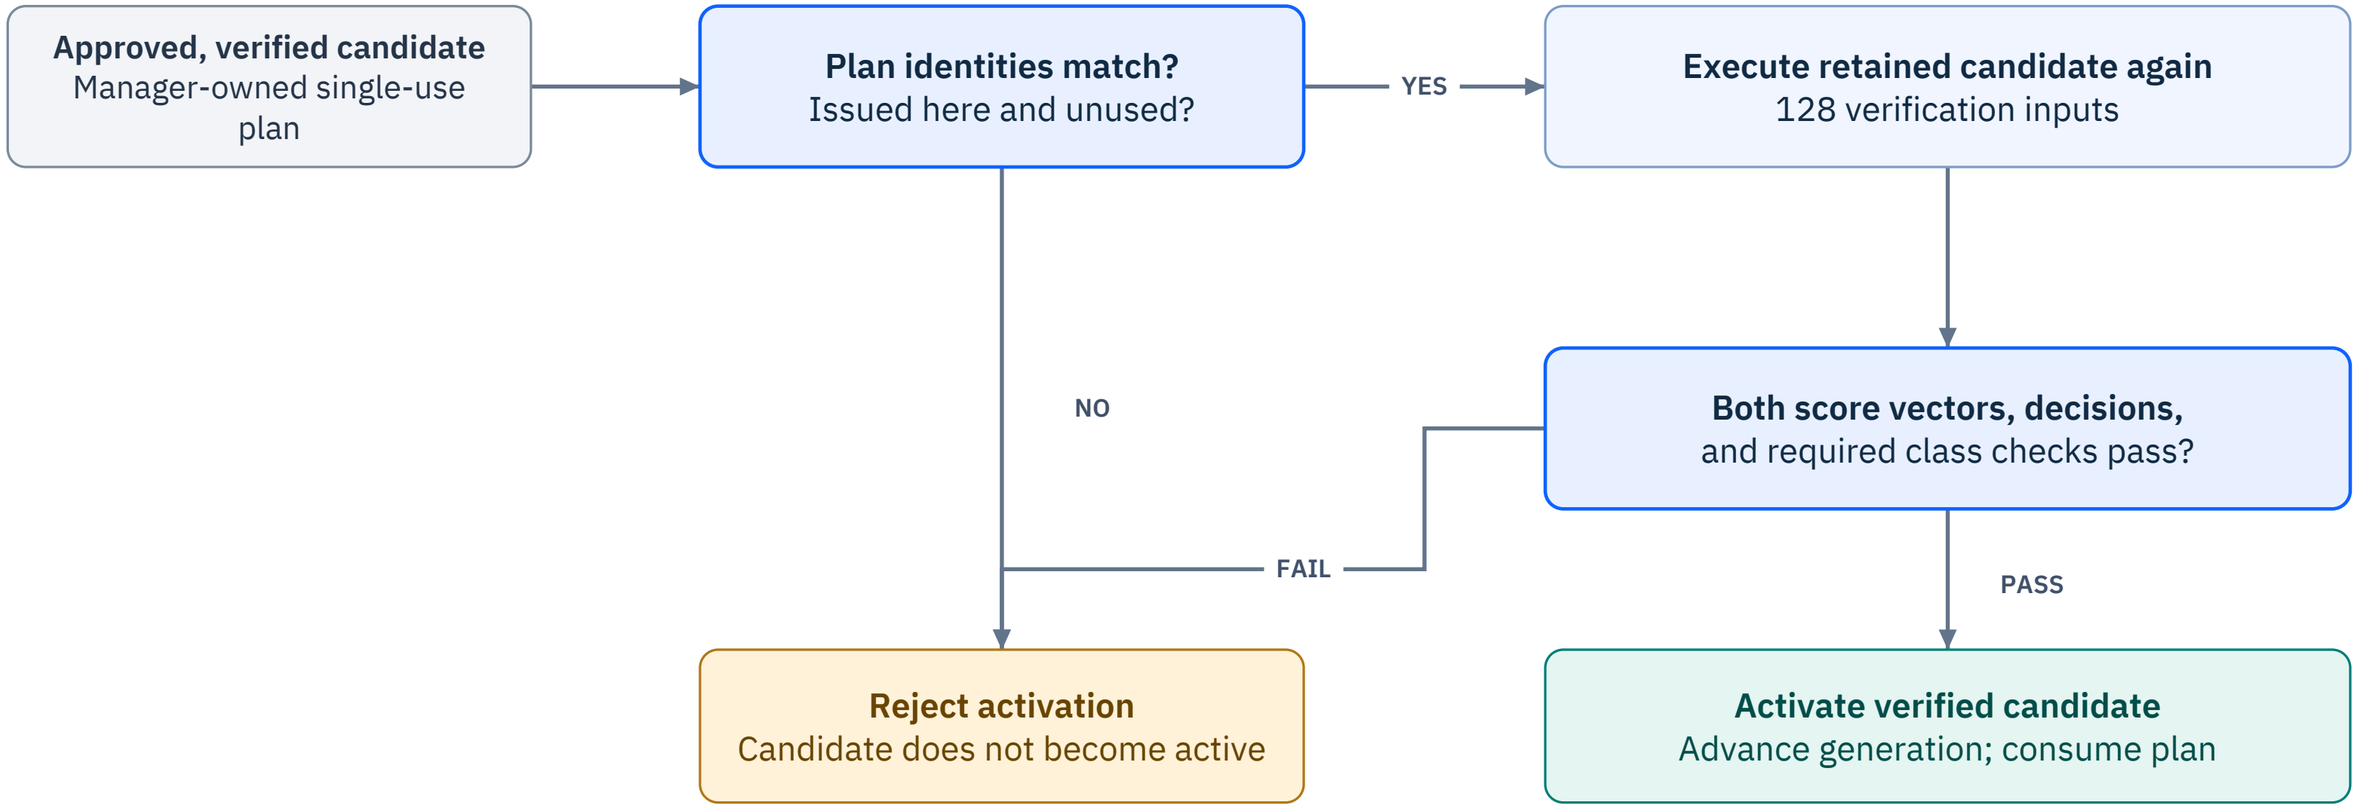

The manager uses its retained candidate and probe inputs, not a caller-supplied report or a cached result.
Activation does not automatically release an existing hold. Separate release approval and full verification still apply.

Data and model lifecycle

Related families stay together before synthetic expansion.

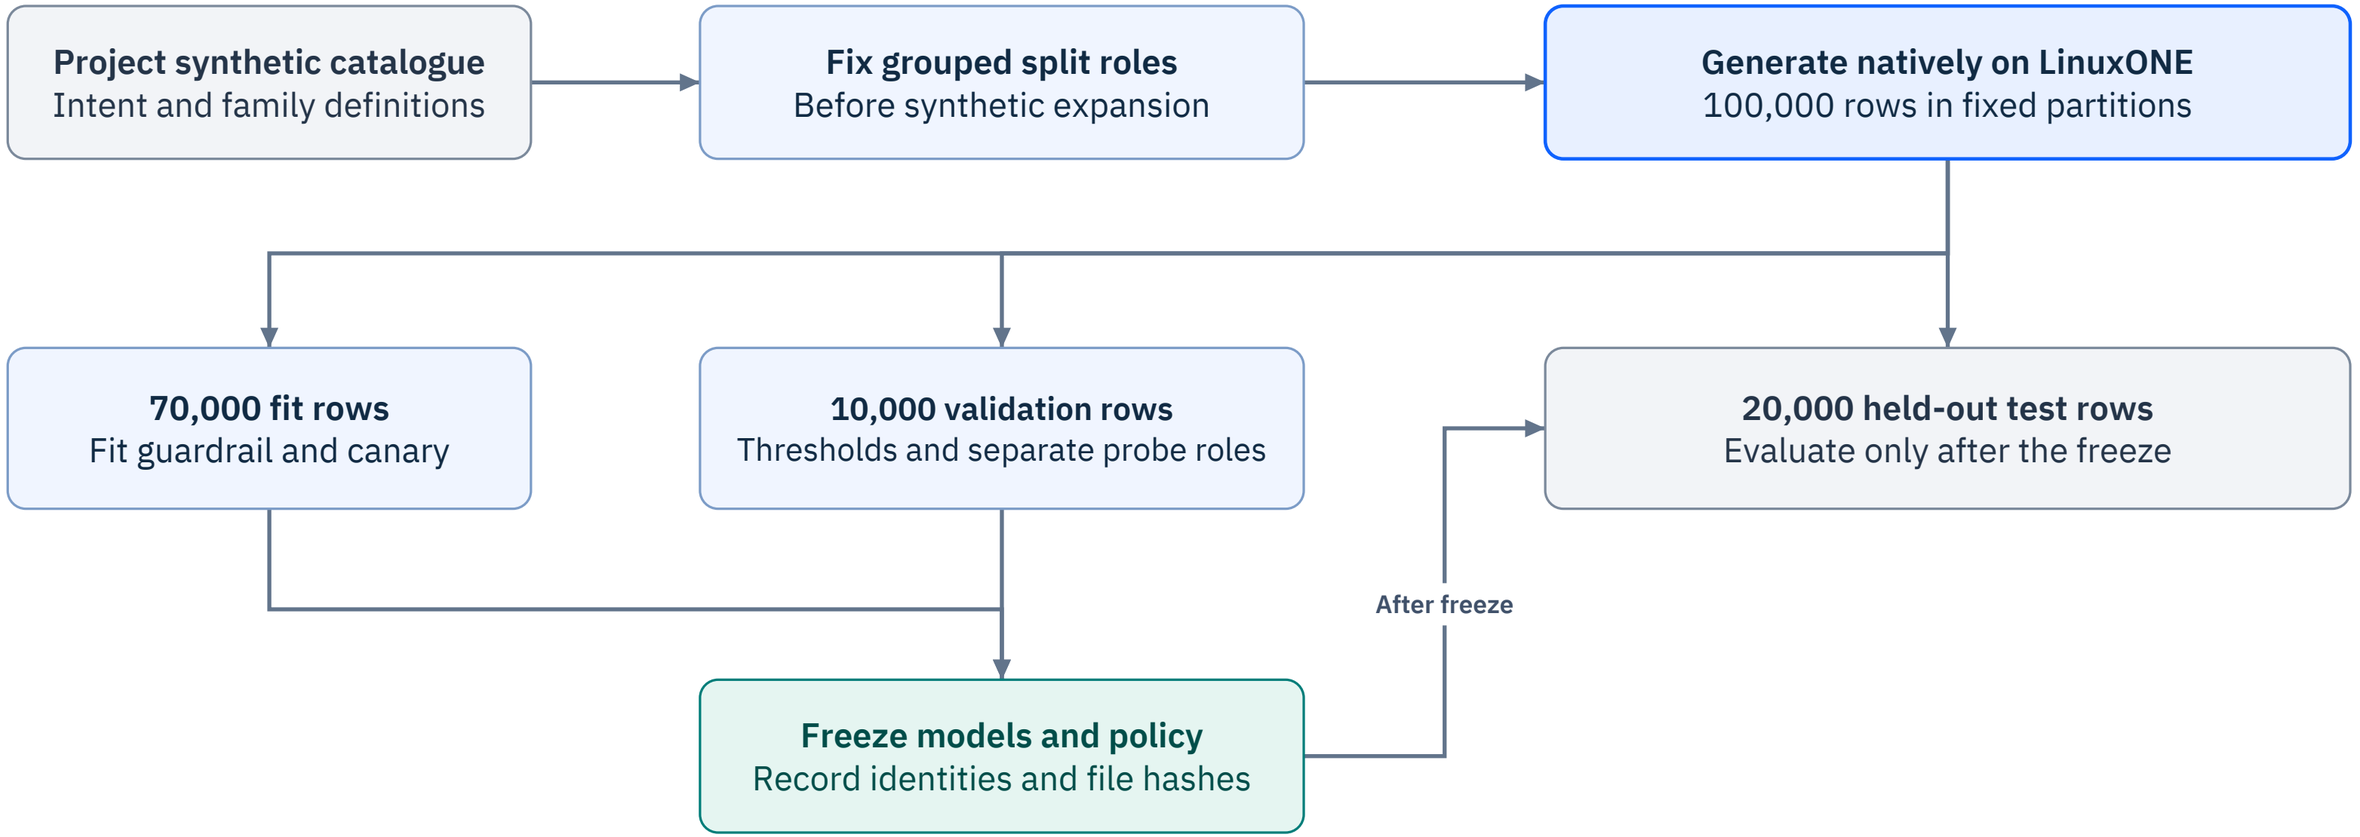

The catalogue contains 100 template families and 91 reviewed semantic groups. Variants are not independent human requests.
The selected model does not use the reserved calibration fallback. Test labels do not select its weights or thresholds.